



# **Sparse Matrix-Vector Multiplication on Throughput-Oriented Processors**

**Nathan Bell and Michael Garland**

NVIDIA Research




# Overview



- **GPUs deliver high SpMV performance**
  - 10+ GFLOP/s on unstructured matrices
  - 140+ GByte/s memory bandwidth
- **No one-size-fits-all approach**
  - Match method to matrix structure
- **Exploit structure when possible**
  - Fast methods for regular portion
  - Robust methods for irregular portion

# Characteristics of SpMV

- **Memory bound**
  - FLOP : Byte ratio is very low
- **Generally irregular & unstructured**
  - Unlike dense matrix operations (BLAS)




# Solving Sparse Linear Systems



- **Iterative methods**



- CG, GMRES, BiCGstab, etc.
- Require 100s or 1000s of SpMV operations



# Finite-Element Methods



- **Discretized on structured or unstructured meshes**
  - **Determines matrix sparsity structure**








# Objectives


- **Expose sufficient parallelism**
  - Develop 1000s of independent threads
- **Minimize execution path divergence**
  - SIMD utilization
- **Minimize memory access divergence**
  - Memory coalescing

# Sparse Matrix Formats



# Compressed Sparse Row (CSR)


- Rows laid out in sequence
- Inconvenient for fine-grained parallelism






# CSR (scalar) kernel

- One thread per row
  - Poor memory coalescing
  - Unaligned memory access



# CSR (vector) kernel


- One SIMD vector or *warp* per row
  - Partial memory coalescing
  - Unaligned memory access



# ELLPACK (ELL)




- **Storage for K nonzeros per row**
  - Pad rows with fewer than K nonzeros
  - Inefficient when row length varies



# Hybrid Format

- ELL handles *typical* entries
- COO handles *exceptional* entries
  - Implemented with segmented reduction




# Exposing Parallelism

- **DIA, ELL & CSR (scalar)**
  - One thread per row
- **CSR (vector)**
  - One warp per row
- **COO**
  - One thread per nonzero



Finer Granularity

# Exposing Parallelism

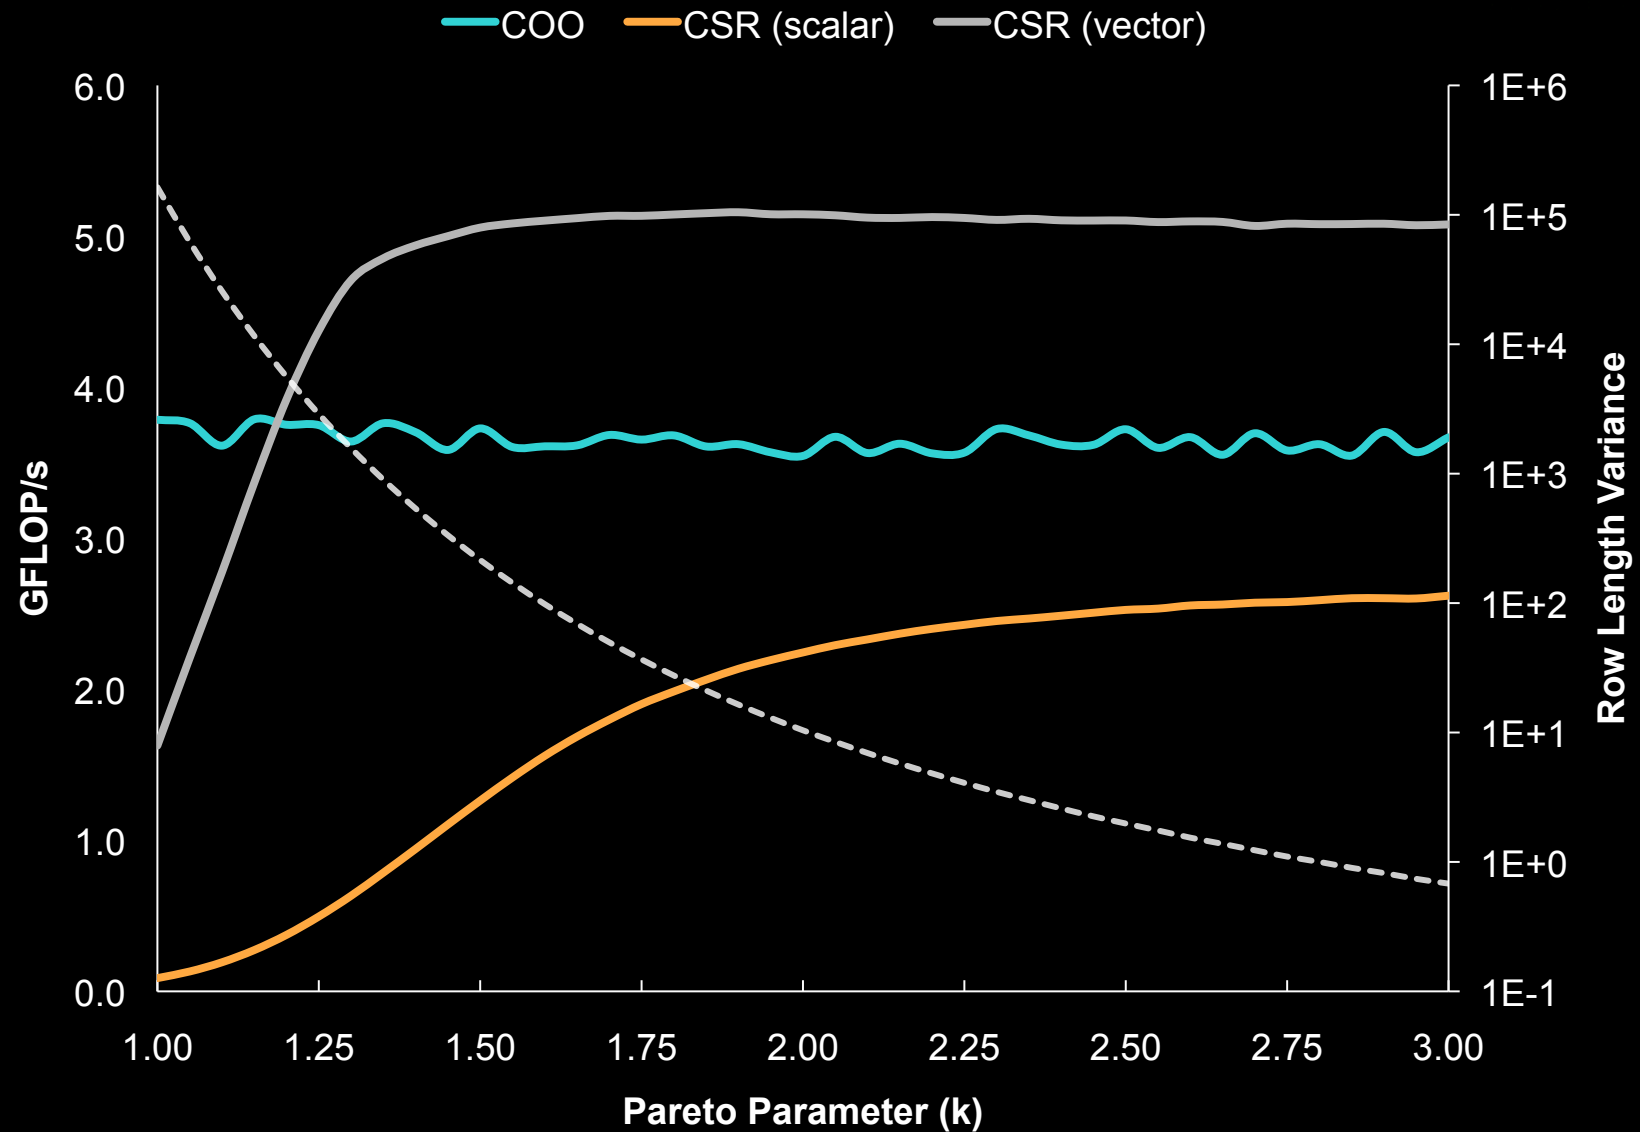




# Execution Divergence

- **Variable row lengths can be problematic**
  - Idle threads in CSR (scalar)
  - Idle processors in CSR (vector)
- **Robust strategies exist**
  - COO is insensitive to row length

# Execution Divergence



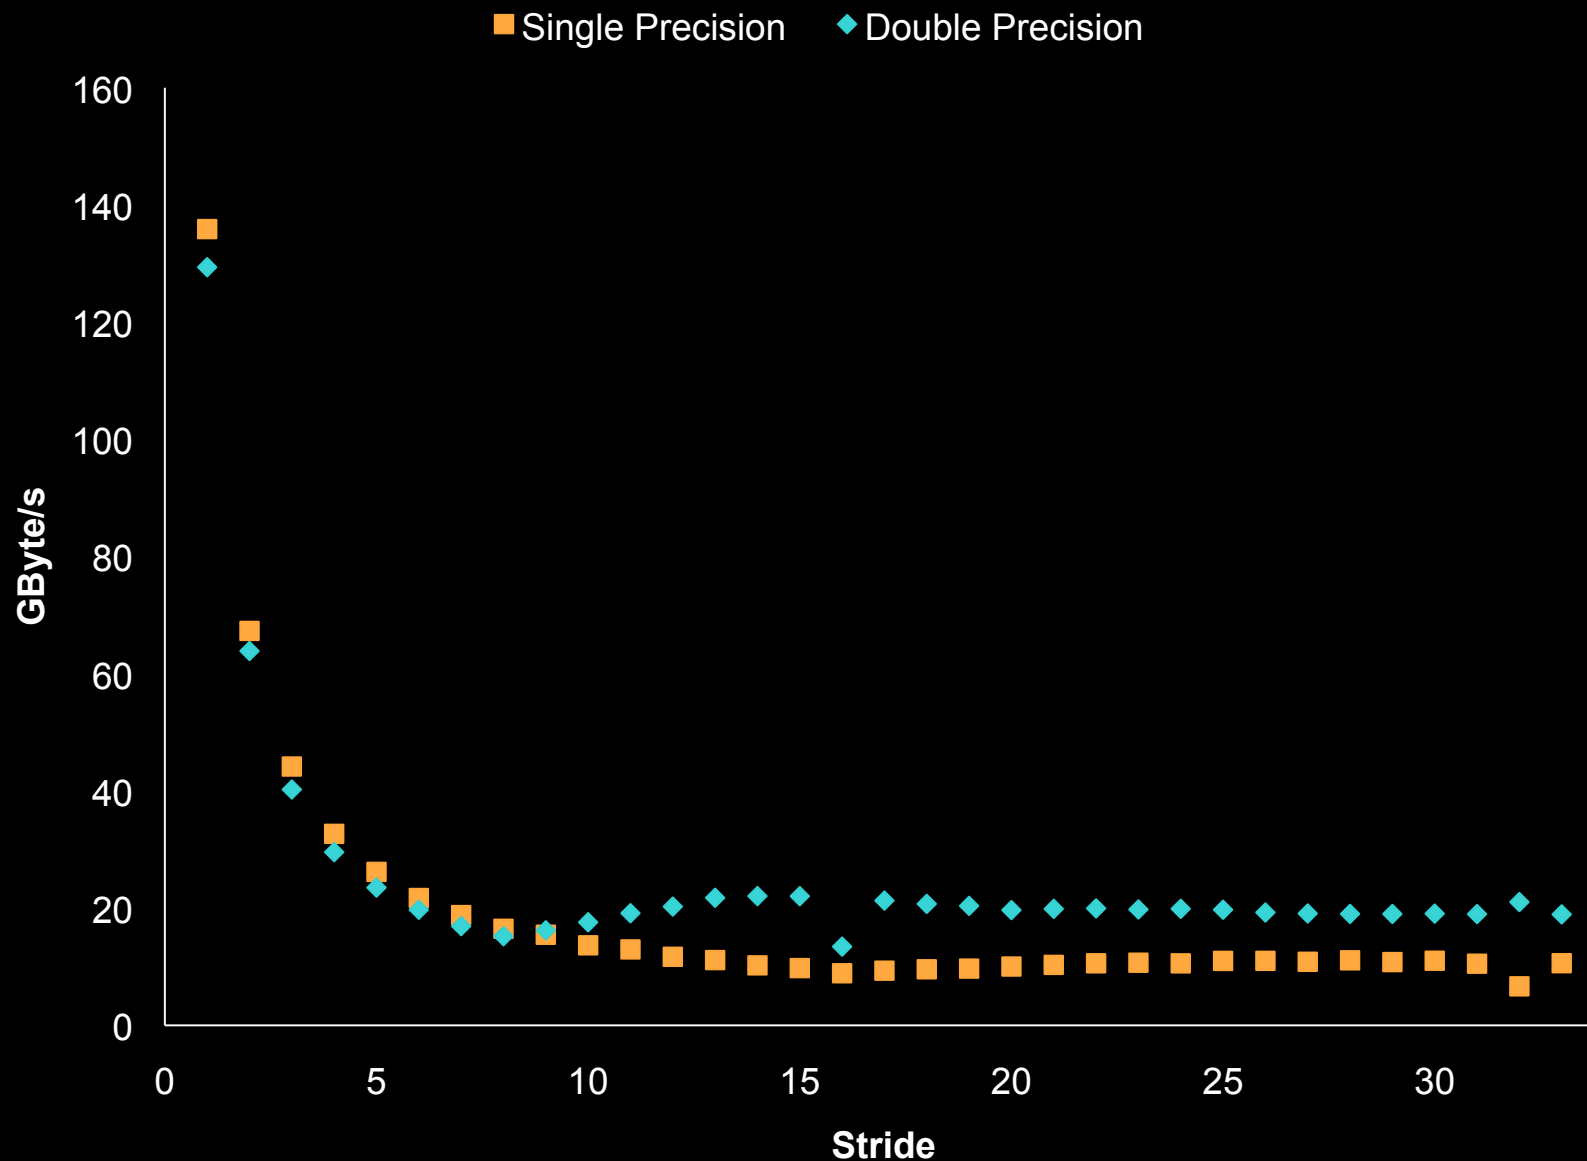




# Memory Access Divergence

- **Uncoalesced memory access is very costly**
  - Sometimes mitigated by cache
- **Misaligned access is suboptimal**
  - Align matrix format to coalescing boundary
- **Access to matrix representation**
  - DIA, ELL and COO are fully coalesced
  - CSR (vector) is partially coalesced
  - CSR (scalar) is seldom coalesced

# Memory Bandwidth (AXPY)




# Performance Results




- **GeForce GTX 285**
  - Peak Memory Bandwidth: 159 GByte/s
  - All results in double precision
  - Source vector accessed through texture cache
- **Structured Matrices**
  - Common stencils on regular grids
- **Unstructured Matrices**
  - Wide variety of applications and sparsity patterns

# Structured Matrices



# Unstructured Matrices



# Performance Comparison



System	Cores	Clock (GHz)	Notes
GTX 285	240	1.5	NVIDIA GeForce GTX 285
Cell	8 (SPEs)	3.2	IBM QS20 Blade (half)
Core i7	4	3.0	Intel Core i7 (Nehalem)

## Sources:

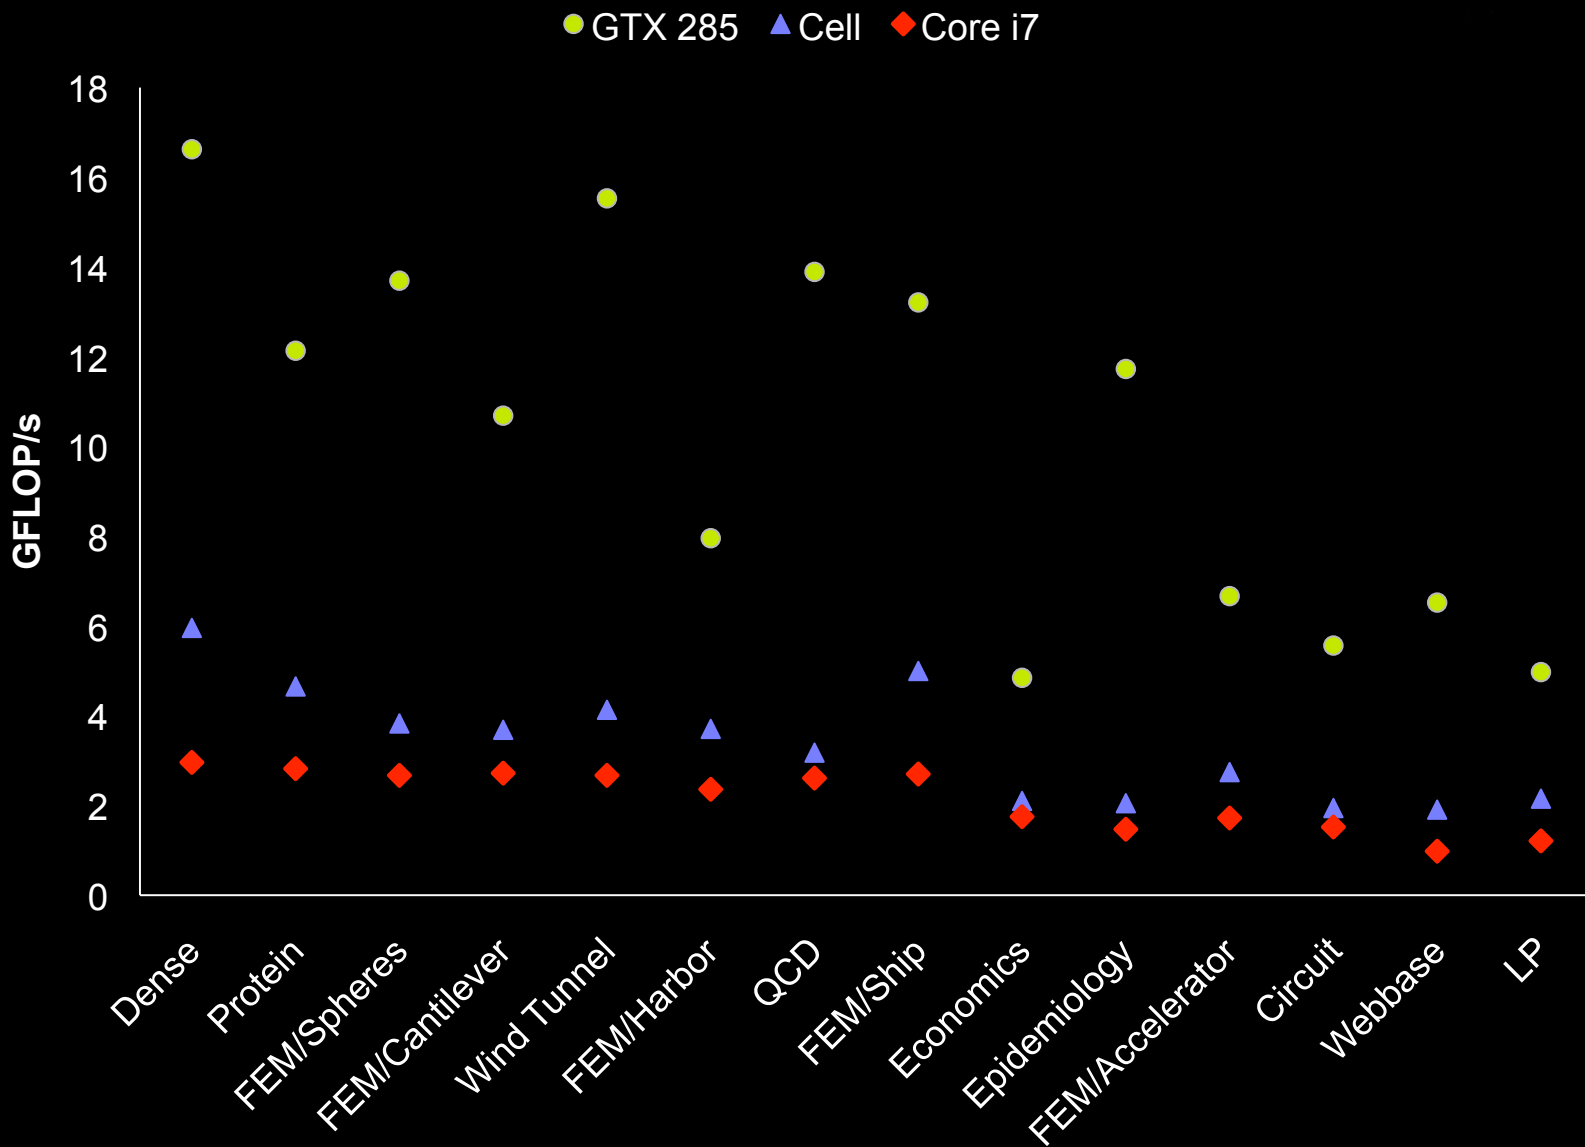
*Implementing Sparse Matrix-Vector Multiplication on Throughput-Oriented Processors*

N. Bell and M. Garland, Proc. Supercomputing '09, November 2009

*Optimization of Sparse Matrix-Vector Multiplication on Emerging Multicore Platforms*

Samuel Williams et al., Supercomputing 2007.

# Performance Comparison



# ELL kernel



```
__global__ void ell_spmv(const int num_rows,          const int num_cols,
                        const int num_cols_per_row, const int stride,
                        const double * Aj,           const double * Ax,
                        const double * x,            double * y)
{
    const int thread_id = blockDim.x * blockIdx.x + threadIdx.x;
    const int grid_size = gridDim.x * blockDim.x;

    for (int row = thread_id; row < num_rows; row += grid_size) {
        double sum = y[row];

        int offset = row;

        for (int n = 0; n < num_cols_per_row; n++) {
            const int col = Aj[offset];

            if (col != -1)
                sum += Ax[offset] * x[col];

            offset += stride;
        }

        y[row] = sum;
    }
}
```





```
#include <cusp/hyb_matrix.h>
#include <cusp/io/matrix_market.h>
#include <cusp/krylov/cg.h>


int main(void)
{
    // create an empty sparse matrix structure (HYB format)
    cusp::hyb_matrix<int, double, cusp::device_memory> A;

    // load a matrix stored in MatrixMarket format
    cusp::io::read_matrix_market_file(A, "5pt_10x10.mtx");

    // allocate storage for solution (x) and right hand side (b)
    cusp::array1d<double, cusp::device_memory> x(A.num_rows, 0);
    cusp::array1d<double, cusp::device_memory> b(A.num_rows, 1);

    // solve linear system with the Conjugate Gradient method
    cusp::krylov::cg(A, x, b);

    return 0;
}
```



<http://cusp-library.googlecode.com>



# Extensions & Optimizations

- **Block formats (register blocking)**
  - Block CSR
  - Block ELL
- **Block vectors**
  - Solve multiple RHS
  - Block Krylov methods
- **Other optimizations**
  - Better CSR (vector)



# Further Reading

## ***Implementing Sparse Matrix-Vector Multiplication on Throughput-Oriented Processors***

N. Bell and M. Garland

Proc. Supercomputing '09, November 2009

## ***Efficient Sparse Matrix-Vector Multiplication on CUDA***

N. Bell and M. Garland

NVIDIA Tech Report NVR-2008-004, December 2008

## ***Optimizing Sparse Matrix-Vector Multiplication on GPUs***

M. M. Baskaran and R. Bordawekar.

IBM Research Report RC24704, IBM, April 2009

## ***Model-driven Autotuning of Sparse Matrix-Vector Multiply on GPUs***

J. W. Choi, A. Singh, and R. Vuduc

Proc. ACM SIGPLAN (PPoPP), January 2010



# Questions?

[nbell@nvidia.com](mailto:nbell@nvidia.com)